

CSCI - 6409 - The Process of Data Science - Fall 2022

</center>

Assignment 3

</center>

Meagan Sinclair

B00737317

Samarth Jariwala

B00899380

1. Data Preparation

1.1 Data Quality Report

```
In [1]: from sklearn import datasets
dataset = datasets.fetch_openml(data_id = 1597, as_frame=True)
```

```
In [2]: df = dataset.frame
df = df.convert_dtypes()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 30 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    V1      284807 non-null  Float64
 1    V2      284807 non-null  Float64
 2    V3      284807 non-null  Float64
 3    V4      284807 non-null  Float64
 4    V5      284807 non-null  Float64
 5    V6      284807 non-null  Float64
 6    V7      284807 non-null  Float64
 7    V8      284807 non-null  Float64
 8    V9      284807 non-null  Float64
 9   V10     284807 non-null  Float64
10   V11     284807 non-null  Float64
11   V12     284807 non-null  Float64
12   V13     284807 non-null  Float64
13   V14     284807 non-null  Float64
14   V15     284807 non-null  Float64
15   V16     284807 non-null  Float64
16   V17     284807 non-null  Float64
17   V18     284807 non-null  Float64
18   V19     284807 non-null  Float64
19   V20     284807 non-null  Float64
20   V21     284807 non-null  Float64
21   V22     284807 non-null  Float64
22   V23     284807 non-null  Float64
```

```
23 V24      284807 non-null  Float64
24 V25      284807 non-null  Float64
25 V26      284807 non-null  Float64
26 V27      284807 non-null  Float64
27 V28      284807 non-null  Float64
28 Amount    284807 non-null  Float64
29 Class     284807 non-null  category
dtypes: Float64(29), category(1)
memory usage: 71.2 MB
```

In [3]: df

Out[3]:

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	0.090794
1	1.191857	0.266151	0.16648	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	-0.166974
2	-1.358354	-1.340163	1.773209	0.37978	-0.503198	1.800499	0.791461	0.247676	-1.514654	0.207643
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	-0.054952
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	0.753074
...
284802	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	4.35617
284803	-0.732789	-0.05508	2.03503	-0.738589	0.868229	1.058415	0.02433	0.294869	0.5848	-0.975926
284804	1.919565	-0.301254	-3.24964	-0.557828	2.630515	3.03126	-0.296827	0.708417	0.432454	-0.484782
284805	-0.24044	0.530483	0.70251	0.689799	-0.377961	0.623708	-0.68618	0.679145	0.392087	-0.399126
284806	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.41465	0.48618	-0.915427

284807 rows × 30 columns

In [67]:

```
classes = df['Class'].to_numpy(dtype='int')
print("Number of class 0: ", len(classes[classes == 0]))
print("Number of class 1: ", len(classes[classes == 1]))
```

```
Number of class 0:  284315
Number of class 1:  492
```

In [4]:

```
# code source: Tutorial [1]
import pandas as pd
import warnings

def build_continuous_features_report(data_df):

    stats = {
        "Count": len,
        "Miss %": lambda df: df.isna().sum() / len(df) * 100,
        "Card.": lambda df: df.nunique(),
        "Min": lambda df: df.min(),
        "1st Qrt.": lambda df: df.quantile(0.25),
        "Mean": lambda df: df.mean(),
        "Median": lambda df: df.median(),
        "3rd Qrt.": lambda df: df.quantile(0.75),
        "Max": lambda df: df.max(),
        "Std. Dev.": lambda df: df.std(),
    }
```

```
contin_feat_names = data_df.select_dtypes("number").columns
continuous_data_df = data_df[contin_feat_names]

report_df = pd.DataFrame(index=contin_feat_names, columns=stats.keys())

for stat_name, fn in stats.items():
    # NOTE: ignore warnings for empty features
    with warnings.catch_warnings():
        warnings.simplefilter("ignore", category=RuntimeWarning)
        report_df[stat_name] = fn(continuous_data_df)

return report_df
```

In [5]: build_continuous_features_report(df)

Out[5]:



	Count	Miss %	Card.	Min	1st Qrt.	Mean	Median	3rd Qrt	Max	Std. Dev.
V1	284807	0.0	275663	-56.407510	-0.920373	1.168375e-15	0.018109	1.315642	2.454930	1.958696
V2	284807	0.0	275663	-72.715728	-0.59855	3.416908e-16	0.065486	0.803724	22.057729	1.651309
V3	284807	0.0	275663	-48.325589	-0.890365	-1.379537e-15	0.179846	1.027196	9.382558	1.516255
V4	284807	0.0	275663	-5.683171	-0.84864	2.074095e-15	-0.019847	0.743341	16.875344	1.415869
V5	284807	0.0	275663	-113.743307	-0.691597	9.604066e-16	-0.054336	0.611926	34.801666	1.380247
V6	284807	0.0	275663	-26.160506	-0.768296	1.487313e-15	-0.274187	0.398565	73.301626	1.332271
V7	284807	0.0	275663	-43.557242	-0.554076	-5.556467e-16	0.040103	0.570436	120.589494	1.237094
V8	284807	0.0	275663	-73.216718	-0.20863	1.205498e-16	0.022358	0.327346	20.007208	1.194353
V9	284807	0.0	275663	-13.434066	-0.643098	-2.406306e-15	-0.051429	0.597139	15.594995	1.098632
V10	284807	0.0	275663	-24.588262	-0.535426	2.238853e-15	-0.092917	0.453923	23.745136	1.088850
V11	284807	0.0	275663	-4.797473	-0.762494	1.673327e-15	-0.032757	0.739593	12.018913	1.020713
V12	284807	0.0	275663	-18.683715	-0.405571	-1.247012e-15	0.140033	0.618238	7.848392	0.999201
V13	284807	0.0	275663	-5.791881	-0.648539	8.190001e-16	-0.013568	0.662505	7.126883	0.995274
V14	284807	0.0	275663	-19.214325	-0.425574	1.207294e-15	0.050601	0.49315	10.526766	0.958596
V15	284807	0.0	275663	-4.498945	-0.582884	4.887456e-15	0.048072	0.648821	8.877742	0.915316
V16	284807	0.0	275663	-14.129855	-0.468037	1.437516e-15	0.066413	0.523296	17.315112	0.876253
V17	284807	0.0	275663	-25.162799	-0.483748	-3.740237e-16	-0.065676	0.399675	9.253526	0.849337
V18	284807	0.0	275663	-9.498746	-0.49885	9.564149e-16	-0.003636	0.500807	5.041069	0.838176
V19	284807	0.0	275663	-7.213527	-0.456299	1.039917e-15	0.003735	0.458949	5.591971	0.814041
V20	284807	0.0	275663	-54.497720	-0.211721	6.407202e-16	-0.062481	0.133041	39.420904	0.770925
V21	284807	0.0	275663	-34.830382	-0.228395	1.656562e-16	-0.029450	0.186377	27.202839	0.734524
V22	284807	0.0	275663	-10.933144	-0.54235	-3.568593e-16	0.006782	0.528554	10.503090	0.725702
V23	284807	0.0	275663	-44.807735	-0.161846	2.610582e-16	-0.011193	0.147642	22.528412	0.624460

	Count	Miss %	Card.	Min	1st Qrt.	Mean	Median	3rd Qrt	Max	Std. Dev.
V24	284807	0.0	275663	-2.836627	-0.354586	4.473066e-15	0.040976	0.439527	4.584549	0.605647
V25	284807	0.0	275663	-10.295397	-0.317145	5.213180e-16	0.016594	0.350716	7.519589	0.521278
V26	284807	0.0	275663	-2.604551	-0.326984	1.683537e-15	-0.052139	0.240952	3.517346	0.482227
V27	284807	0.0	275663	-22.565679	-0.07084	-3.659966e-16	0.001342	0.091045	31.612198	0.403632
V28	284807	0.0	275663	-15.430084	-0.05296	-1.223710e-16	0.011244	0.07828	33.847808	0.330083
Amount	284807	0.0	32767	0.000000	5.6	8.834962e+01	22.000000	77.165	25691.160000	250.120109


```
In [6]: from matplotlib import pyplot as plt
```

```
In [7]: contin_feat_names = df.select_dtypes("number").columns
```


```
In [8]: for column in contin_feat_names:
        range = int(df[column].max() - df[column].min())
        df.hist(column=column, bins=range*2)
        plt.show()
```




V3









V4




V5









V12




V13




V14




V15

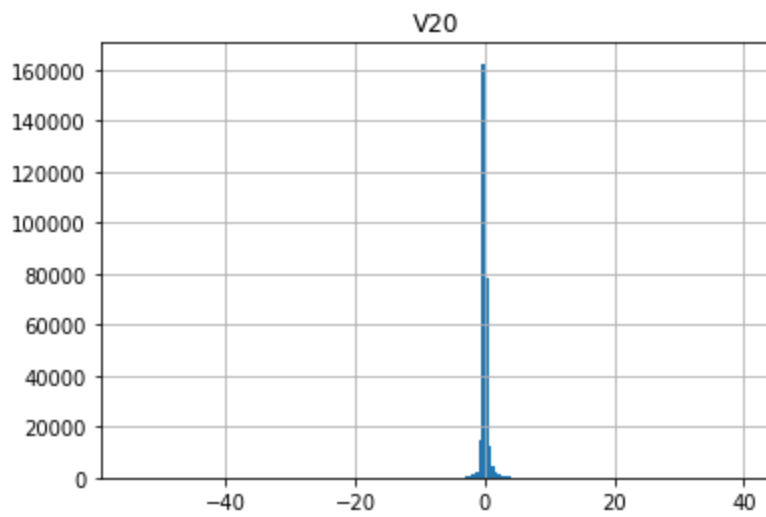







V16




V17









V24






V25











V26

















```
In [9]: for column in contin_feat_names:
        df.boxplot(column)
        plt.show()
```


























1.2 Data Quality Issues and Data Quality Plan

The dataset contains a large number of outliers and removing them would cause even greater class imbalance. Instead, the outlying values will be replaced with a capped min/max value.

1.3 Data Preprocessing

```
In [10]: # Replace the outliers with capped values
# code reference: Tutorial [1]
for column in contin_feat_names:

    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)

    IQR = Q3-Q1
    print(f"IQR = {Q3} - {Q1} = {IQR}")
    outliers_df1 = df[(df[column] < (Q1 - 1.5 * IQR)) ]
    df[column][outliers_df1.index] = (Q1 - 1.5 * IQR)


    outliers_df2 = df[(df[column] > (Q3 + 1.5 * IQR)) ]
    df[column][outliers_df2.index] = (Q3 + 1.5 * IQR)

    tot_outliers = len(outliers_df1) + len(outliers_df2)
    print(column, "Num of outliers: ", tot_outliers, "Percent outliers: ", tot_outliers/len(df[column]))

    df.boxplot(column)
    plt.show()
```


IQR = 1.315641693877865 - -0.920373384390322 = 2.236015078268187

V1 Num of outliers: 7062 Percent outliers: 2.4795738868777804




IQR = 0.8037238712400945 - -0.598549913464916 = 1.4022737847050104

V2 Num of outliers: 13526 Percent outliers: 4.7491810243428




$IQR = 1.027195542465555 - -0.8903648381551406 = 1.9175603806206956$

V3 Num of outliers: 3363 Percent outliers: 1.1807996292225964




$IQR = 0.7433412894685876 - -0.848640116331273 = 1.5919814057998605$

V4 Num of outliers: 11148 Percent outliers: 3.9142296362097846




$IQR = 0.611926439735193 - -0.6915970708876575 = 1.3035235106228504$


V5 Num of outliers: 12295 Percent outliers: 4.316958501722219




$IQR = 0.39856489635610504 - -0.768295608460489 = 1.166860504816594$
V6 Num of outliers: 22965 Percent outliers: 8.063355184388024




$IQR = 0.5704360728775986 - -0.5540758790365226 = 1.1245119519141211$
V7 Num of outliers: 8948 Percent outliers: 3.1417767119487934




$IQR = 0.327345861923449 - -0.2086297440394665 = 0.5359756059629155$
V8 Num of outliers: 24134 Percent outliers: 8.47380857914307




$IQR = 0.5971390302822686 - -0.6430975702665915 = 1.24023660054886$
V9 Num of outliers: 8283 Percent outliers: 2.9082852598426303




$IQR = 0.453923445139507 - -0.5354257264933235 = 0.9893491716328305$
V10 Num of outliers: 9496 Percent outliers: 3.334187713082895




$IQR = 0.739593407321606 - -0.7624941955129775 = 1.5020876028345835$
V11 Num of outliers: 780 Percent outliers: 0.27386967314707855




IQR = $0.618238032946136 - -0.40557148544041355 = 1.0238095183865497$
 V12 Num of outliers: 15348 Percent outliers: 5.3889124916171305




IQR = $0.662504959439974 - -0.6485392991145684 = 1.3110442585545425$
 V13 Num of outliers: 3368 Percent outliers: 1.182555204050462




IQR = $0.493149849218149 - -0.4255740124549935 = 0.9187238616731425$
 V14 Num of outliers: 14149 Percent outliers: 4.96792564789489




$IQR = 0.648820806317158 - -0.582884279157456 = 1.2317050854746139$
V15 Num of outliers: 2894 Percent outliers: 1.0161267103687759




$IQR = 0.523296312475344 - -0.46803676671289796 = 0.991333079188242$
V16 Num of outliers: 8184 Percent outliers: 2.873524878250886




$IQR = 0.3996749826503845 - -0.483748313707048 = 0.8834232963574324$
V17 Num of outliers: 7420 Percent outliers: 2.605273044552978




$IQR = 0.5008067468872159 - -0.498849798665041 = 0.9996565455522569$
 V18 Num of outliers: 7533 Percent outliers: 2.644949035662747




$IQR = 0.458949355762679 - -0.4562989187444475 = 0.9152482745071264$
 V19 Num of outliers: 10205 Percent outliers: 3.5831282236742776




$IQR = 0.1330408409942945 - -0.21172136467424701 = 0.34476220566854154$
 V20 Num of outliers: 27770 Percent outliers: 9.750462593967143




$IQR = 0.1863772033785755 - -0.22839494677851702 = 0.4147721501570925$
V21 Num of outliers: 14497 Percent outliers: 5.090113655914355




$IQR = 0.52855363533339865 - -0.5423503726606616 = 1.0709040079946481$
V22 Num of outliers: 1317 Percent outliers: 0.4624184096598749




$IQR = 0.14764206385605 - -0.16184634501488449 = 0.3094884088709345$
V23 Num of outliers: 18541 Percent outliers: 6.510022576692287




$IQR = 0.439526600168186 - -0.3545861364094985 = 0.7941127365776846$
V24 Num of outliers: 4774 Percent outliers: 1.67622284564635




$IQR = 0.350715562867386 - -0.31714505406527 = 0.667860616932656$
V25 Num of outliers: 5367 Percent outliers: 1.8844340202312442




$IQR = 0.2409521737147555 - -0.3269839258807195 = 0.567936099595475$
V26 Num of outliers: 5596 Percent outliers: 1.9648393473475019




$IQR = 0.09104511968580689 - -0.07083952930446921 = 0.1618846489902761$
V27 Num of outliers: 39163 Percent outliers: 13.750715396742356



$IQR = 0.07827995475782015 - -0.0529597930169809 = 0.13123974777480105$
V28 Num of outliers: 30342 Percent outliers: 10.653530285421356








$IQR = 77.16499999999999 - 5.6 = 71.565$
Amount Num of outliers: 31904 Percent outliers: 11.201971861646658









In [11]:




```
# replot the feature distributions
for column in contin_feat_names:
    range = int(df[column].max() - df[column].min())
    df.hist(column=column, bins=20)
    plt.show()
```














V15




V16




V17




V18









V19






V20









2. Model Training

In [12]:

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import RepeatedStratifiedKFold
```

In [80]:

```
y = df['Class'].to_numpy(dtype=int)
X = df.drop('Class', axis=1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```



```
cv = ShuffleSplit(n_splits=2, test_size=0.2, random_state=0)
```

In [58]:

```
# code source: Tutorial [1]
def plot_learning_curve(
    estimator,
    title,
    X,
    Y,
    axes=None,
    ylim=None,
    cv=None,
    n_jobs=None,
    train_sizes=np.linspace(0.1, 1.0, 5),
):

    _, axes = plt.subplots(1, 3, figsize=(20, 5))

    axes[0].set_title(title)
    if ylim is not None:
        axes[0].set_ylim(*ylim)
    axes[0].set_xlabel("Training examples")
    axes[0].set_ylabel("Score")

    train_sizes, train_scores, test_scores, fit_times, _ = learning_curve(
        estimator,
        X,
        Y,
        cv=cv,
        n_jobs=n_jobs,
        train_sizes=train_sizes,
        return_times=True,
        scoring="accuracy",
    )
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    fit_times_mean = np.mean(fit_times, axis=1)
    fit_times_std = np.std(fit_times, axis=1)

    # Plot learning curve
    axes[0].grid()
    axes[0].fill_between(
        train_sizes,
        train_scores_mean - train_scores_std,
        train_scores_mean + train_scores_std,
        alpha=0.1,
        color="r",
    )
    axes[0].fill_between(
        train_sizes,
        test_scores_mean - test_scores_std,
        test_scores_mean + test_scores_std,
        alpha=0.1,
        color="g",
    )
    axes[0].plot(
        train_sizes, train_scores_mean, "o-", color="r", label="Training score"
    )
    axes[0].plot(
        train_sizes, test_scores_mean, "o-", color="g", label="Cross-validation score"
    )
    axes[0].legend(loc="best")
```

```

# Plot n_samples vs fit_times
axes[1].grid()
axes[1].plot(train_sizes, fit_times_mean, "o-")
axes[1].fill_between(
    train_sizes,
    fit_times_mean - fit_times_std,
    fit_times_mean + fit_times_std,
    alpha=0.1,
)
axes[1].set_xlabel("Training examples")
axes[1].set_ylabel("fit_times")
axes[1].set_title("Scalability of the model")

# Plot fit_time vs score
fit_time_argsort = fit_times_mean.argsort()
fit_time_sorted = fit_times_mean[fit_time_argsort]
test_scores_mean_sorted = test_scores_mean[fit_time_argsort]
test_scores_std_sorted = test_scores_std[fit_time_argsort]
axes[2].grid()
axes[2].plot(fit_time_sorted, test_scores_mean_sorted, "o-")
axes[2].fill_between(
    fit_time_sorted,
    test_scores_mean_sorted - test_scores_std_sorted,
    test_scores_mean_sorted + test_scores_std_sorted,
    alpha=0.1,
)
axes[2].set_xlabel("fit_times")
axes[2].set_ylabel("Score")
axes[2].set_title("Performance of the model")
return plt

```

2.1 Strong Learner: Gaussian Naive Bayes

source: [7]

Gaussian Naive Bayes is used when the data is continuous because it computes the probabilities of the likelihoods. The data should have Gaussian distribution ideally, our visualization plots show Gaussian-like distribution. Also, can be accurate with just a few data points, which is ideal as we have limited number of Class 1 data. The data requires strong independence between the features. Our dataset is majority transformed to hide the original feature information but feature independence is a reasonable assumption in this case as previous data scientists processed the features into a usable set.

```

In [14]: from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import GridSearchCV, RepeatedStratifiedKFold
from sklearn.model_selection import learning_curve
import matplotlib.pyplot as plt

```

Hyperparameter tuning

```

In [81]: # code source : [2]

import warnings
warnings.filterwarnings('ignore')

model = GaussianNB()

params_NB = {'var_smoothing': np.logspace(0, -9, num=10)}
gs_NB = GridSearchCV(estimator=model,
                    param_grid=params_NB,

```

```

cv=cv,
verbose=1,
scoring='accuracy')
gs_NB.fit(X_train, y_train)

```

Fitting 2 folds for each of 10 candidates, totalling 20 fits

```

Out[81]: GridSearchCV(cv=ShuffleSplit(n_splits=2, random_state=0, test_size=0.2, train_size=None),
    estimator=GaussianNB(),
    param_grid={'var_smoothing': array([1.e+00, 1.e-01, 1.e-02, 1.e-03, 1.e-04,
    1.e-05, 1.e-06, 1.e-07,
    1.e-08, 1.e-09])},
    scoring='accuracy', verbose=1)

```

```

In [82]: gs_NB.best_params_

```

```

Out[82]: {'var_smoothing': 0.0001}

```

Build and train model

```

In [83]: model = GaussianNB(var_smoothing=gs_NB.best_params_['var_smoothing'])
    model.fit(X_train, y_train);

```

```

In [84]: y_pred = model.predict(X_test)

```

```

In [85]: confusion_matrix(y_test, y_pred)

```

```

Out[85]: array([[93792,   46],
    [   25,  124]])

```

```

In [86]: print(classification_report(y_test, y_pred))

```




	precision	recall	f1-score	support
0	1.00	1.00	1.00	93838
1	0.73	0.83	0.78	149
accuracy			1.00	93987
macro avg	0.86	0.92	0.89	93987
weighted avg	1.00	1.00	1.00	93987

Plot and analyze learning curve

```

In [87]: estimator = GaussianNB(var_smoothing=gs_NB.best_params_['var_smoothing'])
    p = plot_learning_curve(estimator, "GNB", X, y, cv=cv)
    plt.show()

```



The learning curve shows that model achieves good training score as the increase in the training examples. The model achieves best score with approx. 150000 training examples and after that the model is overfitting with decrease in cross-validation score.

The model fit_time increases with training examples and also the performance of the model.

2.2 Bagging Model

Sklearn's ensembled BaggingClassifier uses DecisionTreeClassifier by default as the base estimator and ensembles multiple estimators. The DecisionTreeClassifier is a good option for our data as classification is our goal and this model is low cost. Both numerical and categorical data is supported without requiring preprocessing so our data meets the requirements. Ensembling these models with bagging may provide better results than a single decision tree.

Source: [3]

```
In [22]: from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import GridSearchCV, RepeatedStratifiedKFold
```

Hyperparameter tuning

```
In [89]: param_grid = {
        'n_estimators' : [2, 5, 7 ]
    }

    clf = GridSearchCV(BaggingClassifier(),
                      param_grid,
                      cv=cv,
                      scoring = 'accuracy')

    clf.fit(X_train, y_train)

    print(clf.best_params_)

{'n_estimators': 7}
```

Build and train model

```
In [94]: model = BaggingClassifier(n_estimators=7)

        # weight the class 1 samples higher to compensate for the sample imbalance
        #sample_weight = np.ones(shape=(len(y_train),))
        #sample_weight[y_train == 1] = 100
        # Note: the weighting value was selected with trial and error

        model = model.fit(X_train, y_train, )
```

```
In [95]: y_pred = model.predict(X_test)
```

```
In [96]: confusion_matrix(y_test, y_pred)
```


```
Out[96]: array([[93823,    15],
               [   28,   121]])
```

```
In [97]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	93838
1	0.89	0.81	0.85	149
accuracy			1.00	93987
macro avg	0.94	0.91	0.92	93987
weighted avg	1.00	1.00	1.00	93987

Learning curve

```
In [98]: estimator = BaggingClassifier(n_estimators=7)
p = plot_learning_curve(estimator, "BC", X, y, cv=cv)
plt.show()
```



The learning curve for bagging classifier shows that the training score is good throughout the increment of training examples. On the other side, the cross-validation score increases. The fit timing of the model also increases with training examples. The performance of the model never decreases with an increase of training examples.

2.3 Boost Model

AdaBoostClassifier was chosen as it also uses DecisionTreeClassifier as the base estimator, but implements a Boost algorithm for ensemble. This will be interesting to compare with the bagging model. The same data requirements as above.

Source: [4]

```
In [31]: from sklearn.ensemble import AdaBoostClassifier
```

Tune hyperparameters

```
In [32]: param_grid = {
          'n_estimators' : [5, 10, 15 ]
        }

clf = GridSearchCV(AdaBoostClassifier(),
                  param_grid, scoring = 'accuracy')

clf.fit(X_train, y_train)

clf.best_params_
```

```
Out[32]: {'n_estimators': 10}
```

Build and train model

```
In [53]: model_ADA = AdaBoostClassifier(n_estimators=10)

        model_ADA = model_ADA.fit(X_train, y_train, )
```

```
In [54]: y_pred = model_ADA.predict(X_test)
```

```
In [55]: confusion_matrix(y_test, y_pred)
```


```
Out[55]: array([[93806,    32],
               [   40,   109]])
```

```
In [56]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	93838
1	0.77	0.73	0.75	149
accuracy			1.00	93987
macro avg	0.89	0.87	0.88	93987
weighted avg	1.00	1.00	1.00	93987

Learning Curve

```
In [99]: estimator = AdaBoostClassifier(n_estimators=10)
        p = plot_learning_curve(estimator, "ABC", X, y, cv=cv)
        plt.show()
```



The training score has been seen decreasing with increase of training examples. The cross validation score is first increasing with some training examples but after that it also has been seen decreasing.

The performance of the model has been seen gradually decreasing after a certain amount of time model has taken to train.

3. Model Comparison

3.1 Variance and Bias Analysis

The Strong learner model has overfitting in the data which might arised from variance and bias trade-off tension between the errors.

The variance has been decreased by the bagging classifier model, therefore throughout the increase in training examples, the score has been highest and same.

In the boosting model, it is seen that the training score decreases but the cross validation score increases with increase with training examples upto a point. Then both becomes stable, that is, the model is biased to a particular kind of solution.

The bagging classifier model is suitable for fraud detection because it has the best accuracy and macro-avg f1 score compared to the rest of models. Also, it has good training score throughout and the model is trained better with more training examples as the cross-validation score, scalability and performance of the model also increased.

Source: [5]

Macro average f1 calculates the F1 separated by class but not using weights for the aggregation. This results in a bigger penalisation when the model does not perform well with the minority classes which is exactly what we want when there is imbalance.

3.2 Imbalance Classification Analysis

Source: [6] The problem with the fraud detection dataset is that the majority of the data is non-fraudulent transactions (Class 0) and contains very few fraudulent (Class 1). The overall accuracy may appear to be high if the model tries to classify most or all data points as Class 0 but fails to complete the goal of identifying Class 1. Imbalanced data like this must be handled appropriately to ensure the business problem is solved.

The model which can handle the imbalanced data best is the ADA boosting model. As ADA boost builds ensemble of weak learners by adjusting weight of misclassified data during each iteration, higher weight is given to the minority class as they will be misclassified more often if the model tries to become biased towards Class 0. To achieve similar accuracy with the other ensemble model, sample weights were required to be used, otherwise Class 1 accuracy was very poor. The strong learner model

3.3 Run Time Analysis


The run time analysis curves with 1 worker thread are plotted in section 2 along with the learning curves. The fastest model in this case is gaussian naive bayes model.

In [100...

```
from sklearn.model_selection import ShuffleSplit

title = "Learning Curves for Guassian Naive Bayes"
n_jobs = 4


estimator = GaussianNB(var_smoothing=gs_NB.best_params_['var_smoothing'])
plot_learning_curve(estimator, title, X_train, y_train, cv=cv, n_jobs=n_jobs)
plt.show()
```



In [101...

```
title = "Learning Curves for Bagging Classifier"
```


```
estimator = BaggingClassifier(n_estimators=15)
plot_learning_curve(estimator, title, X_train, y_train, cv=cv, n_jobs=n_jobs)
plt.show()
```



In [102...

```
title = "Learning Curves for AdaBoost Classifier"
```

```
estimator = AdaBoostClassifier(n_estimators=10)
plot_learning_curve(estimator, title, X_train, y_train, cv=cv, n_jobs=n_jobs)
plt.show()
```



The fastest model in the case of 4 worker threads is the Gaussian Bayes classifier.

References

- [1] In-class Tutorial \ [2] <https://stackoverflow.com/questions/39828535/how-to-tune-gaussiannb> \ [3] <https://scikit-learn.org/stable/modules/tree.html#tree> \ [4] <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html> \ [5] <https://datascience.stackexchange.com/questions/65839/macro-average-and-weighted-average-meaning-in-classification-report> \ [6] <https://towardsdatascience.com/https-medium-com-abrown004-how-to-ease-the-pain-of-working-with-imbalanced-data-a7f7601f18ba> \ [7] <https://iq.opengenus.org/gaussian-naive-bayes/>